

Written Assignment 2 : Summer-2 (2023)
Priyanka Bojja : 700739528

1. Considering the join operation between relation r and s ($r \bowtie_{\theta} s$), θ is $r \cdot A = s \cdot B$ with
- Relation r contains 20,000 tuples and has 10 tuples per block.
 - Relation s contains 2,000 tuples and has 10 tuples per block.
 - there are 17 buffer blocks available in memory (M).

(1.1) Using Block Nested Loop Join:

$$\text{cost} = \left(\left\lceil \frac{b_r}{M-1} \right\rceil * b_s \right) + b_r$$

$$b_r = \frac{20,000}{10} = 2000$$

$$b_s = \frac{2000}{10} = 200$$

$$\text{cost} = \left(\left\lceil \frac{2000}{17-1} \right\rceil * 200 \right) + 2000$$

$$\text{cost} = (125 * 200) + 2000$$

Cost = 27,000 block transfers

(1.2) Using Merge Join:

$$\begin{aligned} B_g &= b_r * \left(2 \left\lceil \log_{M-1} \left(\frac{b_r}{M} \right) \right\rceil + 2 \right) + b_s * \left(2 \left\lceil \log_{M-1} \left(\frac{b_s}{M} \right) \right\rceil + 2 \right) \\ &= 2000 * \left(2 \left\lceil \log_{17-1} \left(\frac{2000}{17} \right) \right\rceil + 2 \right) + 200 * \left(2 \left\lceil \log_{17-1} \left(\frac{200}{17} \right) \right\rceil + 2 \right) \end{aligned}$$

$$= 2000 * (2 \lceil 1.719 \rceil + 2) + 200 * (2 \lceil 0.889 \rceil + 2)$$

$$= 2000 * (2(2)+2) + 200(2(1)+2)$$

$$= 2000 * 6 + 200 * 4 = 12,000 + 800 = 12,800$$

$$\text{Cost} = B_s + b_r + b_s$$

$$= 12,800 + 2000 + 200$$

= 15,000 Block transfers

(1.3) Hash join (r \bowtie_0 s): Non recursive partition:

$$\text{Cost} = 3(B_s + b_r)$$

$$= 3(2000 + 200)$$

= 6,600 block transfers

(1.4) Hash Join (r \bowtie_0 s): Recursive partition:

$$\text{Cost} = 2(B_s + b_r) \lceil \log_{M-1}(B_s) - 1 \rceil + b_r + b_s$$

where ($B_s \leq b_r$)

$$= 2(2000 + 200) \lceil \log_{16}(200) - 1 \rceil + 2000 + 200$$

$$= 4,400 \lceil 0.91 \rceil + 2000 + 200$$

= 6,600 Block transfers

(1.5) If there is the infinity of memory, then the total minimum cost will be 6,600 Block transfers i.e.; using hash join (No-recursive) because it is independent of memory for the cost calculation.

(2) SQL Query to find the instructor's name who teaches 'Advanced Database' in Fall Semester, 2021

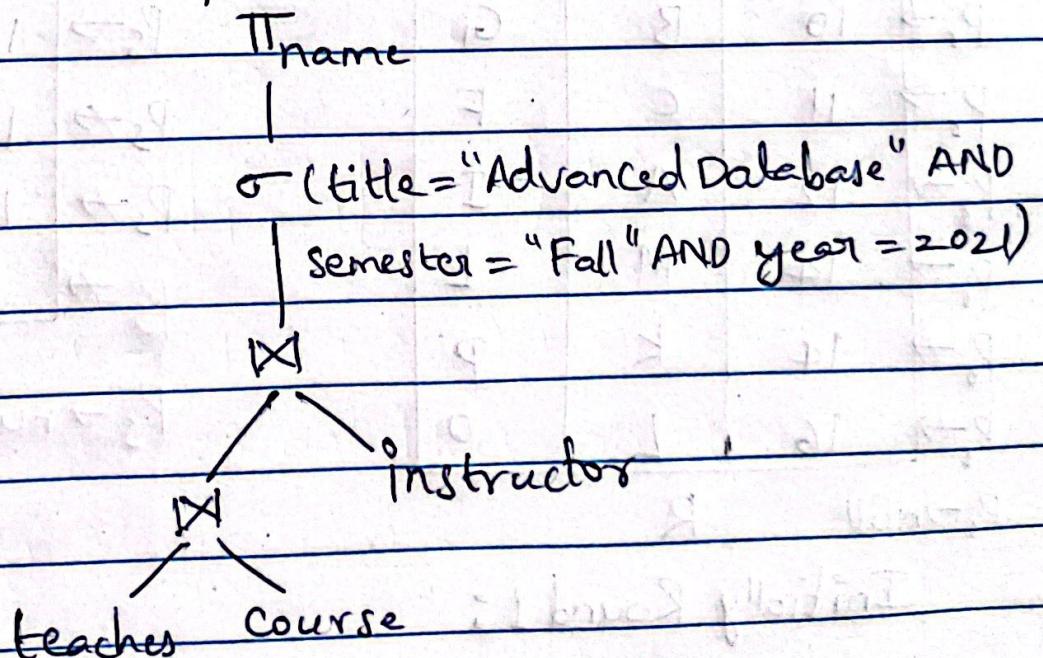
Query : SQL Command :

```
SELECT name FROM instructor, teaches, course  
WHERE teaches.course_id = course.course_id AND  
instructor.id = teaches.id AND title = "Advanced  
Database" AND semester = "Fall" AND year = 2021;
```

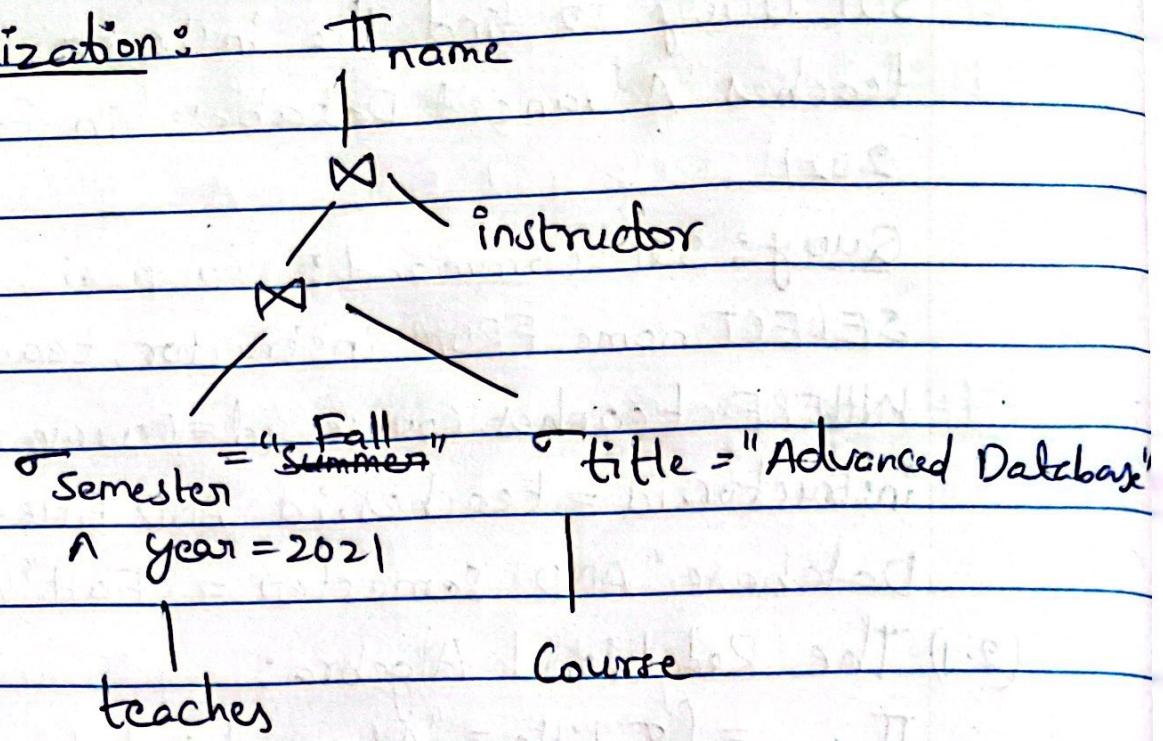
(2.1) The Relational Algebra :

$$\Pi_{\text{name}} = (\sigma_{\text{title} = \text{"Advanced Database"} \text{ AND } \text{Semester} = \text{"Fall"} \text{ AND } \text{year} = 2021}$$
$$(\text{Instructor} \bowtie (\text{teaches} \bowtie \text{course}))$$

(2.2) Equivalent Expression :



optimization:



③ The Merge join with the Samples relation R and S

	A ₁	A ₂	A ₃
P _r ↗	10	A	L
P _r ↗	10	B	G
P _r ↗	11	C	F
P _r ↗	12	A	I
P _r ↗	14	M	L
P _r ↗	14	K	P
P _r ↗	16	L	O

R
 $P_r \rightarrow \text{null}$

	A ₁	A ₂
P _s ↗	10	30
P _s ↗	12	40
P _s ↗	12	50
P _s ↗	15	20
P _s ↗	16	10

S
 $P_s \rightarrow \text{null}$

Initially Round 1:

Where $P_r = (10, A, L)$ $P_s = (10, 30)$

$t_s = (10, 30)$ $S_s = \{(10, 30)\}$

updated $P_s = (12, 40)$

done: false $t'_s = (12, 40)$

Round #	t_s	S_s	t_r
1	(12, 40)	{(10, 30)}	(11, C, F)
2	(15, 20)	{(12, 40), (12, 50)}	(14, M, L)
3	(16, 10)	{(15, 20)}	(16, L, 0)
4	null	{(16, 10)}	null

updated done : true

$$t_r = (10, A, L) \otimes t_s(10, 30) \Rightarrow$$

$$\text{updated } P_r = (10, B, G)$$

$$t_r = (10, B, G) \otimes t_s(10, 30) \Rightarrow$$

$$\text{updated } P_r = (11, C, F)$$

$$t_r = (11, C, F)$$

Result			
A1	A2	A3	A4
10	A	L	30
10	B	G	30
12	A	I	40
12	A	I	50
16	L	O	10

$$\text{Round 2: } P_r = (11, C, F) \quad P_s = (12, 40)$$

$$t_s = (12, 40)$$

$$S_s = \{(12, 40)\}$$

$$t_s' = (12, 50)$$

$$\text{updated } P_r = (12, A, I)$$

$$S_s = S_s \cup t_s' = \{(12, 40), (12, 50)\}, P_s = (12, 50)$$

$$\text{updated } P_r = (14, M, L), t_r = (14, M, L), P_s = (15, 20)$$

$$\text{Round 3: } P_r = (14, M, L), P_s = (15, 20), t_s = (15, 20), S_s = \{(15, 20)\}$$

$$t_s' = (16, 10)$$

$$\text{updated } P_r = (14, K, P)$$

$$\text{again updated } P_r = (16, L, O), t_r = (16, L, O)$$

$$t_s = (16, 10), S_s = \{(16, 10)\}$$

$$\text{Round 4: } P_r = (16, L, O), P_s = (16, 10), t_s = (16, 10)$$

$$S_s = \{(16, 10)\}, P_s \rightarrow \text{null}$$

$$P_r \rightarrow \text{null}, t_r \rightarrow \text{null}$$

	S1			S2		
(4)	T ₁	T ₂	T ₃	T ₁	T ₂	T ₃
Read(z)				Read(z)		
write(y)				write(y)		
	Read(x) write(x) Read(z)			Read(x) commit		
		Read(x) write(x)			Read(x) write(x)	
		write(z) Commit			Read(z) write(z)	
Read(x)					Commit	Read(x)
Commit			Read(y) Commit			write(x)
						Read(Y) Commit

4.1 In schedule S1 the conflicts In schedule S2 the conflicts are as follows : are as follows :

T₁ Read(z) with T₂ write(z) T₁ Read(z) with T₂ write(z)

T₁ write(y) with T₃ Read(Y) T₁ ^{write} Read(Y) with T₃ Read(Y)

T₂ Read(x) with T₃ write(x) T₁ read(x) with T₂ write(x)

T₂ write(x) with T₃ Read(x) T₁ read(x) with T₃ write(x)

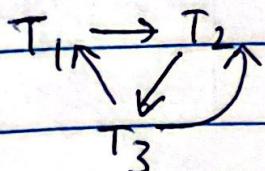
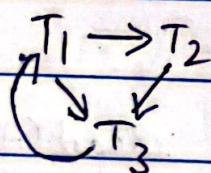
T₃ write(x) with T₁ Read(x) T₂ read(x) with T₃ write(x)

T₂ write(x) with T₃ read(x)

T₃ read(x) with T₂ write(x)

T₃ write(x) with T₂ Read(x)

T₃ read(y) with T₁ write(y)



By comparing the set of conflicts, we can see that S_1 and S_2 have different sets of conflict. Therefore S_1 is not conflict equivalent to S_2 .

4.2 View equivalent where

	S_1	X	Y	Z		S_2	X	Y	Z
I IR		T_2	T_3	T_1		I IR	T_1	T_3	T_1
II P/C		$T_2 \rightarrow T_3$ $T_3 \rightarrow T_1$	$T_1 \rightarrow T_3$	$T_1 \rightarrow T_2$		II P/C	T_2/T_3	T_1/T_3	-
III LUP		T_3	T_1	T_2		III LUP	T_3	T_1	T_2

Therefore S_1 and S_2 do not produce the same values for all data items. So, S_1 is not view equivalent to S_2 .

5. Schedule S_3

	T1	T2
	Lock-X(A) Read(A) $A = A + 100$	Lock-S(B) Read(B) Lock-SCA)
	write(A) unlock(A) Lock-X(B)	Read(A) unlock(B) commit unlock(A) Display(A+B)
	Read(B) $B = B - 100$ write(B) commit unlock(B)	

5.1) In schedule S3, we can identify the following dependencies and locks:

- T_1 acquires a lock(X) on A before reading and updating.
- T_2 acquires a lock(S) on B before reading.
- T_2 acquires a shared lock on A after releasing the lock on B.

The deadlock does not occur in Schedule S3.

Although there is a dependency between T_1 and T_2 , the locking mechanism allows for compatibility between two transactions.

Therefore no deadlock in S3.

5.2) The schedule S3 already guarantees serializability with T_1 followed by T_2 .

T_1 holds an exclusive lock on A and performs its operations before releasing the lock.

T_2 acquires shared lock on B and A and commits. No modification is necessary to guarantee serializability. T_1 followed by T_2 .